# Alex Turing

San Francisco, CA | alex.turing@email.com | linkedin.com/in/alexturing | github.com/alexturing
AI / ML Engineer with 8+ Years of Experience

## Professional Summary

Results-driven AI/ML Engineer with 8 years of experience designing, training, and deploying scalable machine learning models and artificial intelligence solutions. Proven expertise in natural language processing (NLP), computer vision, and predictive analytics. Adept at bridging the gap between complex data science concepts and robust production environments through comprehensive MLOps practices. Strong track record of leading cross-functional teams to deliver enterprise-grade AI products.

## Core Competencies & Technologies

**Machine Learning & AI:** Deep Learning, NLP, Large Language Models (LLMs), Computer Vision, Predictive Modeling, Reinforcement Learning
**Frameworks & Libraries:** PyTorch, TensorFlow, Keras, Scikit-Learn, LangChain, Hugging Face Transformers, Pandas, NumPy
**MLOps & Cloud:** AWS (SageMaker, EC2, S3), GCP (Vertex AI), Docker, Kubernetes, MLflow, Apache Airflow, CI/CD pipelines
**Programming Languages:** Python, SQL, C++, R, Bash

## Professional Experience

**TechGlobal Solutions** | San Francisco, CA
*Senior AI/ML Engineer* (Jan 2021 – Present)

- Lead the design and deployment of enterprise-scale generative AI applications, integrating custom fine-tuned LLMs into core customer service products, which reduced response times by 40%.
- Architect end-to-end MLOps pipelines using AWS SageMaker, Docker, and Kubernetes, ensuring seamless model training, evaluation, and deployment with 99.9% uptime.
- Mentor a team of 4 junior data scientists and machine learning engineers, driving best practices in code quality, version control, and reproducible experiments.
- Develop an automated model monitoring system with MLflow to detect data drift, triggering automatic retraining cycles when accuracy drops below 95%.

**DataVision Analytics** | Seattle, WA
*Machine Learning Engineer* (June 2018 – Dec 2020)

- Engineered a deep learning-based recommendation engine for an e-commerce platform using PyTorch, increasing user engagement metrics by 22% within the first quarter of deployment.
- Optimized data ingestion pipelines processing over 10TB of structured and unstructured data daily using Apache Spark and Airflow.
- Collaborated closely with product managers and backend engineering teams to translate business requirements into actionable machine learning infrastructure.
- Implemented RESTful APIs using FastAPI to serve inference requests with sub-50ms latency.

## Notable Projects

**Autonomous Customer Support Agent:** Built an autonomous conversational agent utilizing LangChain and OpenAI APIs, integrated with internal knowledge bases via vector databases (Pinecone) to achieve a 75% first-contact resolution rate.

**Predictive Maintenance for IoT:** Developed time-series forecasting models (LSTM) to predict equipment failure in manufacturing, saving an estimated $1.2M annually in downtime costs.

## Education

**Master of Science in Computer Science (Specialization in Artificial Intelligence)**
University of Technology | Graduated: May 2018

**Bachelor of Science in Software Engineering**
State University | Graduated: May 2016