Jordan Patel

New York, NY | jordan.patel@email.com | linkedin.com/in/jordanpatel-ai | github.com/jpatel-ml

Staff AI Engineer | 8 Years of Experience

Professional Summary

Visionary Staff AI Engineer with 8 years of end-to-end experience specializing in Computer Vision, Edge AI, and applied Generative AI. Proven ability to bridge hardware and software by optimizing complex neural networks for low-latency inference environments. Highly skilled in building robust Retrieval-Augmented Generation (RAG) pipelines and scalable MLOps architectures. Passionate about transforming academic research into high-impact, production-ready enterprise systems.

Technical Skills

AI & Machine Learning: Computer Vision (Object Detection, Segmentation), Applied GenAI, RAG, Transfer Learning, XGBoost
Frameworks & Tools: PyTorch, TensorFlow, TensorRT, Triton Inference Server, OpenCV, LlamaIndex, vLLM, Qdrant (Vector DB)
Cloud & MLOps: Azure Machine Learning, AWS EC2/S3, Docker, Kubernetes, GitHub Actions, Terraform
Languages: Python, C++, CUDA, SQL, Go

Professional Experience

Nova Visionary Tech | New York, NY
Staff AI Engineer (Feb 2023 – Present)

- Spearheaded the development of a low-latency enterprise RAG platform, utilizing vLLM and LlamaIndex to query internal proprietary documents, accelerating employee workflow efficiency by 35%.
- Optimized multi-modal foundation models for edge deployment using NVIDIA TensorRT, reducing inference latency by 60% without compromising accuracy.
- Led a cross-functional team of 6 engineers to build scalable infrastructure using Kubernetes and Triton Inference Server, managing up to 10M+ daily API requests.

Nexus Autonomous Solutions | Boston, MA
Senior Machine Learning Engineer (Aug 2020 – Jan 2023)

- Developed real-time object detection and semantic segmentation pipelines using YOLO and PyTorch for autonomous warehouse robotics.
- Implemented fully automated CI/CD pipelines for ML models, enabling seamless training, validation, and over-the-air (OTA) updates to edge devices.
- Created synthetic data generation pipelines to handle edge-case scenarios, boosting the model robustness and mAP (mean Average Precision) by 15%.

DataCorp Analytics | Chicago, IL
Machine Learning Engineer (June 2018 – July 2020)

- Built predictive models using XGBoost and LightGBM to analyze customer churn, resulting in targeted retention campaigns that saved $2M annually.
- Designed extensive data preprocessing and feature engineering pipelines in Pandas and SQL to handle highly imbalanced datasets.
- Collaborated with software engineering teams to wrap legacy Python ML scripts into robust REST APIs using Flask and Docker.

Education

Bachelor of Science in Electrical Engineering & Computer Science
Northeastern University | Graduated: May 2018